Money Moves the Pen:
Link Prediction in Congress Bill Co-Sponsorship Networks Using Political Donor Network Information

Yi Zhong, yizhon@stanford.edu
Eddie Chen, yc4@stanford.edu

December 9, 2018

Abstract

We combined the campaign contribution records and congressional bill co-sponsorship data to construct a tri-partite graph, in order to model the money politics in the US Congress. We found that the US Congressional Politics is indeed a small world with collaboration resembling academic collaborations. More importantly, we modeled the bill co-sponsorship prediction as a link prediction problem, using attributes learned from campaign contribution networks. Result shows that using campaign networks is a good way to predict future in-office collaboration between legislators, especially with a supervised, decision-tree model.

1 Introduction

Political collaboration is an important part of legislative life, and congress bill cosponsorships provide a rich source of information about the social network between legislators [4], and serving as a proxy to understand legislators' "connectedness" and collaboration graph. Moreover, according to Mark Twain, "we have the best government that money can buy" - money and politics have already been intertwined. In this project, we applied social network analysis tools on political donation networks and congress bill cosponsorship networks, and framed our research problem as a link prediction task on congress bill cosponsorship networks using political campaign donation records for the US (Congress and Presidential Campaigns) with its network characteristics. We modeled and presented graph characteristics of the two political networks, and showed investigation results of link prediction using various supervised learning techniques for this project. We then compared models' performance to a naive baseline to come up with evaluations.

2 Related Work

While there is a substantial amount of literature in studying the congress networks and the link prediction problem, no prior work exists on combining congress bill sponsorship network with campaign money networks and apply link prediction algorithms to the combined graph. Below, we review some of the state-of-the-art papers on this topic.

Fowler in his *Connecting the Congress: A Study of Cosponsorship Networks* [4] mapped the cosponsorship networks of all 280,000 pieces of legislation proposed in the U.S. House and Senate from 1973 to 2004, and reported out

several interesting statistics about the resulting networks. He further proposed a measure of "connectedness" by looking at the frequency of cosponsorship and the number of cosponsors on each bill to estimate the social "distance" between legislators, and used "connectedness" as a proxy for legislative influence. While the paper has done an excellent exploration around America's political networks from an unique angle, it leaves more to be desired. He treated all the links as unweighted in the paper; but as he himself pointed out, some cosponsors are probably more important than others. Another downside of this paper is that it ignores the temporal aspects of the co-sponsorship network. It only looks at each Congress in this isolation without change over time. Yet, given the dynamic nature of bill writing and co-sponsoring, a study of how co-sponsorship network forms and evolves (link prediction) can reveal insight on how one can work on being more connected and more influential in legislative outcomes.

Dominguez [3] examined the makeup of the Democratic and Republican party coalitions by analyzing the contribution patterns of elite donors (defined as all individual donors who gave over $200 to one of the two major political party committees in the 2003-2004 election cycle) and their network patterns. He showed that both parties are similar in their degree of centralization, with the party committees being the most central actors in the network.

Both political donor networks (in the form of Super PACs) and Congress bill cosponsorships have been studied by students in CS224W before. In **Co-Authorship Networks in Congress** [9], the authors looked at the impact of co-sponsorship on legislation success. They used the network characteristics to predict future co-sponsorship, via **Supervised Random Walks and Triads**. While this past project provided a lot good ideas, it lacked a discussion on using other machine learning models, which can potentially yield good results, especially when the network is dense and Supervised

Random Walks are known to perform well on sparse graphs.

In **Super-PAC Donor Networks** [7], the author studied individual donors and their contributions to specifically Super PACs, a new form of political action committees that can raise unlimited sum of money for campaigns. The authored looked at community structures and other network characteristics for insights on partisan polarization. The author did not show how networks evolve over time, as donations might swing back and forth depending on which party is in power, and the change in donor demographics.

For Link Prediction, **Liben-Nowell et al.** [6] discusses similarity based methods for link prediction, and focuses on using network intrinsic traits to predict its future edge evolution. It explores a wide array of similarity measures (such as Common Neighbors, Jaccard's Index, Adamic/Adar, and Katz clustering, etc.) and compares their prediction accuracy among themselves and against a random predictor as baseline. However, the definition of node similarity is a nontrivial challenge, and it's likely that different networks would require different definition of node "proximity". Moreover, a lot of the similarity measures assume that the link itself indicated similarity between two nodes, which may or may not be true.

Al Hasan et al. [1] models the link prediction problem as a supervised learning task, identifies a set of proximity, aggregated, and topological feartures, and applies common supervised algorithms like decision tree, SVM, k-NN, etc. It evaluates these different models using metrics such as accuracy, recall, F-values,etc. However, the paper skips entirely on hyperparameter fine-tuning for the model of choice (SVM), likely because of the high accuracy it was able to achieve with its well selected features. Moreover, the paper landed luckily on a balanced data set with roughly equal numbers of positive and negative labels; yet for most tight-knit communities that resemble a "small

world", labels are likely skewed, and we need to pay close attention to data selection for training, and probably considering techniques like down-sampling.

Backstrom et al [2] proposed a new hybrid solution in tackling link prediction problem. It uses Random Walks with Restarts (basically Personalized PageRank with one node) as a way to explore the network structure. It then looks at the node and edge attribute data to bias the random walk so that it will more often visit nodes to which creates edges in the future, in a direct and principled way and not to overfit. Yet, the paper only considers node and edge attribute data, and posits that such intrinsic structure likely reflects exogenous characteristics as well (in the Facebook friending example, ones location/closeness in network to other people reflects the likelihood of people partying together and therefore adding on Facebook together). Its unclear to me if this holds true in political networks, especially in the Congress, when politicians come from vastly different places all over the country and from different ideological allegiances as well.

The related work reviewed above provided a lot great ideas, and the most obvious one is perhaps to marry political donor network, congress bill co-sponsorship network, and link prediction together. Below, we present 3 creative contributions our project explored.

1. **Model networks**: Both sections above deal with bipartite graphs, and our subject of study is a tripartitite graph consisting of donors, politicians and bills. Unlike Fowler [4], we modeled networks as **undirected**, which is more compatible with existing link prediction literature.

2. **Incorporate edge weights in supervised learning models**: Because our network of interest here carries a lot important information such as **donation amount**, We extended the abovementioned work by incorporating edge weights.


Figure 1: Illustration of the Congress Political Network

3. **Compare features learned from graph to candidate information**: We ran models based on features learned from candidates' campaign donation networks versus models based on candidates' party and home state information, to understand what information is more predictive.

3 Method

3.1 Problem Statement

Our project is made up of two parts: graph modeling, and link prediction. For graph modeling, we aim to construct a tripartite graph of political committees, legislators (we will ignore those failed to get elected to office), and the bills those legislators worked together on. A sample graph can be found in Figure 1. With the graph constructed, we provide a set of statistics and descriptions of the graph structure (including their one-mode projections, for both bills-legislators and committees-legislators subgraphs). After that, we construct a link prediction problem by dividing graph into different years of congress, and select the suitable years for model training and evaluation. Lastly, we report our learnings from the entire exercise.

Practically, we hope our research can quanti-

3

tatively answer the question: **Does donation in election affect collaboration in office?** Given a graph of congressional politicians and their campaign donations, we want to predict who will co-sponsor bills together as a form of political collaboration. Here we focus on congressional bill (including all resolutions, bills, and amendments) co-sponsorship because co-sponsorship is an observable signal, and tells us intuitively how much support one bill has and therefore how much clout the politician behind the bill has.

This problem has obvious utility: it would be useful to keep the electorate informed of their elected representatives' political collaborations and alliances in the Congress, as well as indicate any changes or tendencies in politicians' stances on various issues. Moreover, for politicians, this information can be used as a guidance to seek more targeted co-sponsors, sparing them from wasting precious political capital and writing countless (and potentially spammy and ineffective) "Dear Colleagues" letters.

3.2 Data Preparation

In this project, we used the campaign financial data provided by the Federal Election Committee[1] from 1981 to 2016 (the 97th to the 114th Congress, including House, Senate and Presidential races). The bill co-authorship data is obtained from the Government Publishing Office's website for the same period. [2]. We have hosted our code repository at https://github.com/yzhong94/cs224w-project/.

A large part of effort to date has been devoted to data cleaning. In particular, we have to join the campaign financial data, which uses FEC's Candidaite IDs, with the Congress bill data, which uses its own ID system, on legislators. Our approach is to first join by legislator's capitalized full first name + last name

+ state (abbreviation), which leaves out more than 100 legislators not being matched between the two datasets. A close examination reveals that some legislators go by nicknames in signing bills but have full legal names on the campaign financial records - for instance, Tom Lewis is actually Thomas Lewis, and James Cooper going by Jim; this is made worse by irregular nicknames too, such as Richard "Doc" Hastings. Some legislators go by their middle names, such as Raymond Eugene Green going by Gene Green, David Adam Smith by Adam Smith. To combat this, we first filtered out all the candidates not matched by the method above (capitalized full first name + last name + state abbreviation), and joined them together by capitalized full last name + state abbreviation, because people are very unlikely to have nicknames for their last names. After that, we created an Excel check function to alert us if we have one NodeID (from the bill data) maps to multiple different Candidate IDs (note: it's possible for a person to have two different Candidate IDs, which happens when this person ran for House first, then Senate later). We then manually inspected the flagged rows and kicked out false positives. In the end, we were able to find 1813 legislators/candidates from both the campaign financial record and the bill co-sponsorship data from 1981 to 2016.

3.3 Network Construction

We constructed a tripartite network: committees, legislators, bills. A committee can be a PAC, SuperPAC or party committee. A legislator is an elected official in the Congress, which can be a senator or a representative. A bill is a draft of legislation that has been sponsored by a legislator and co-sponsored by others. To describe the whole graph, we include all years of data first (from 1981 to 2016), and then look at one term's data for an individual graph.

Between committees and legislators, an undirected link is added if a committee donates to a

---

[1]FEC Website data: https://www.fec.gov/data/advanced/?tab=bulk-data

[2]https://www.govinfo.gov/bulkdata/BILLS/115

legislator (we do not allow for multiple edges between two nodes). This way, we ended up having 1813 candidates with donations across the years. We aggregated the donation amounts between candidates and committees by year and preserved the sum as edge weights, to simplify multi-edges. Between legislators and bills, an undirected link is added first if candidates appear on the bill as either an author or a cosponsor.

3.4   Link Prediction

After constructing the graph, we applied **supervised learning link prediction** on the bill cosponsorship part of the tripartitite graph, **between politicians and bills**.

Formally, let $G(V, E)$ be our entire tripartitite graph with node set V and edge set E, covering periods from t_{start} to t_{end}. Let t_0 be an arbitrary time period between t_{start} to t_{end}. Our training and test graph pair will be $\{G_{t_0}, G_{t_1}\}$: we will train based on network characteristics found in the subgraph $G_{t_0}^{campaign}(V_{legislators}, V_{committees})$, and use that to predict edge formations among **the candidate nodes** in $G_{t_1}^{co-sponsor}(V_{legislators}, V_{bills})$ where $V \in \{G_{t_0}, V_{legislators}\}$ (which is $V_{legislators,t_1}$). We can then repeat the process, run the best model on a new pair of graphs (for different years) as validation, and report the metrics for final evaluation.

We frame our link prediction problem as follows: predict the link between legislators, where a link exists if two legislators cosponored a bill together, for a specific Congress term.

We have constructed four link prediction models:

- Naive baseline predictor: a baseline model based on graph density

- Legislator only predictor: a baseline model based on candidate attributes (party, state)

- Campaign only predictor: prediction model using features generated from campaign

graph network attributes

- FIXME Candidate and campaign predictor: prediction model using features from both campaign graph network attributes and candidate attributes

3.4.1   Naive baseline predictor

We define our **naive baseline predictor** as follows: given a pair of nodes v_1, v_2, we will always predict there will be an edge between these two pairs, i.e. as a complete graph. This is computed for $G_{co-sponsor}$. That is,

$$Accuracy_{NaiveBaseline} = \frac{2||E||}{||V||(||V|| - 1)}$$

Using the 100th Congress (1987 - 1988) as the training set and the 101th Congress (1989-1990) as the test set. The baseline accuracy is calculated as $96,052/138,075 = 0.695$ per above.

3.4.2   Legislator only predictor

We define a second baseline using **legislator attributes only. Features for link prediction all come from information about the legislators**. There are two features: IsFromSameState, IsFromSameParty.

Formally, for every $E(V_{legislator,i}, V_{legislator,j}$ in $G_{co-sponsor}$, IsFromSameState is 1 if $V_{legislator,i}$ and $V_{legislator,j}$ are from the same state, and 0 other wise, likewise for IsFromSameParty. Features are then used in machine learning models for link prediction.

3.4.3   Campaign only predictor

Features for link prediction all come from the campaign network *prior to* the Congress going into session, which is a bipartite network (legislators and committees), where a link between a committee node and a legislator node exists if the committee donates money to the legislator. We want to see if immediate donation has an effect on collaboration,

hence we use campaign data two years before to predict the cosponsorship network during a congressional term.

For example, if we are predicting cosponsorship in the 100th congress (1987 - 1988), we would use campaign data from 1985 to 1986, in order to construct the features.

We tried two types of feature construction:

- Supervised feature learning using network structure

- Unsupervised feature learning using node embedding from node2vec random walks

For generated features, we constructed features from the campaign subgraph solely. Features include:

- Common Neighbors, Union of Neighbors, Jaccard Index

- Degree Difference in a pair of legislator nodes

- Contribution Amount (sum and absolute difference)

- Clustering Co-efficient (sum, absolute difference, mean)

- Degree Centrality difference

- Shortest Distance between two legislator nodes

- Spectral Clusters from Clauset-Newman-Moore greedy modularity maximization

Before feeding all the features engineered above into our machine learning models, we conducted feature selection as well, using Scikit-Learn's implementation of F-statistic in an ANOVA test.

For features construction using node embeddings, we used shallow encoding, $ENC(v) = Zv$, and node2vec algorithm for random walks. We used the example implementation from[5] with the following parameter:

- $p = 1$, $q = 2$ for BFS-like walks

- walk length $= 80$, number of walks $= 10$

Using node embeddings learned from random walks, we computed features using the following aggregation function:

- Hadamard: $f(z_i, z_j) = g(z_i * z_j)$

- Sum: $f(z_i, z_j) = g(z_i + z_j)$

- Average: $f(z_i, z_j) = g(\frac{z_i + z_j}{2})$

- Distance: $f(z_i, z_j) = g(\|z_i - z_j\|_2)$

Before feeding all the features learned above into our machine learning models, we conducted feature selection as well, again using Scikit-Learn's implementation of F-statistic in an ANOVA test to select the top 20 percentile.

For models to predict link, we have tried two algorithms: logistic regression and decision tree. For logistic regression, we used scikit-learn's [8] default implementation with -1,1 notation for labels and $L2$ regularization. The optimization problem formulation is as

$$min_{w,c}\frac{1}{2}w^T w + C\sum_{i=1}^{n}\log(e^{(-y_i(X_i^T w + c)) + 1})$$

A decision tree is a tree where each node represents a feature, each branch represents a decision/rule and each leaf represents a classification in our case. We used Scikit Learn's default implementation which uses Gini Index as the metric [8]. Specifically, we define $C = \{-1, 1\}$ as our target class, E as the set of records where $(E_1, E_2, ...E_k)$ will be the splits induced on E. We aim to decrease the impurity measure, which is measured by the Gini Index (Perfectly classified, Gini Index would be zero).

Let p_j be the fraction of records in E of class c_j,

$$p_j = \frac{|t \in E : t[C] = c_j|}{|E|}$$

Then, as we have 2 classes in our case,

$$Gini(E) = 1 - \sum_{j=1}^{2} p_j^2$$

3.4.4 Evaluation Method

We used accuracy as our main success measure:

$$Accuracy = \frac{NumberOfCorrectPredictions}{TotalNumberOfPredictionsMade}$$

4 Results and Findings

4.1 Network Description

The basic stats of the tripartitie graph are included below:

- Legislator count: 1,919 (1813 of which are found in campaign financial network)

- Bill count: 221,726

- Committee count: 14,326

- Edges between legislators and bills: 3,086,039

- Edges between committees and candidates: 911,965

- Overall tripartite graph node count: 237,971, and edge count: 3,998,004

In order to understand clustering coefficients of each parts of the graph, we have divided it into "bill" and "campaign" subgraphs by applying **graph folding** to the respective bipartite graphs (legislators-bills and legislators-committees, both folding to the legislator nodes). As a result, the bill subgraph Clustering coefficient: **0.821170** while the campaign subgraph has a clustering coefficient of **0.988841** - both are very high numbers, indicating that both subgraphs represent a very small and tightly connected world. We are dealing with very dense graphs.

The bill subgraph's highest degree is 11,316 for any legislators (connecting to bills), while for bills it is 433 (so a top bill can garner 433 co-sponsors...for reference, the entire US House


Figure 2: Overall Tripartite Graph Degree Distribution on log-log scale


Figure 3: Degree Distribution of Legislator Nodes in the Bill Subgraph on a linear scale


Figure 4: Degree Distribution of Legislator Nodes in the Campaign Subgraph on a linear scale

has 435 seats). Similarly, the campaign subgraph's highest is 2,093 for legislators connecting to political committees, while the highest for any committees to connect with candidates is 1,669 - this could be the Democratic and/or Republican Party Committee that provides support to all their party's candidates.

In addition, we have plotted the degree distributions of the overall tripartite graph in 2.

It's perhaps more informative to look at candidates' degree distribution in the context of each subgraphs as well; so we have plotted degree distributions for both subgraphs, for legislator nodes in 3 and 4.

Moreover, we have applied role detection to both subgraphs, with the "average legislator" as the comparison baseline and looking at the same three features as HW2: the degree of node c, the number of edges in the egonet of v, and the number of edges connects v's egonet and the rest of the graph. The "average legislator" is defined as a hypothetical node with average values of

Figure 5: Roles in Bill Subgraph

Figure 6: Roles in Campaign Subgraph

the 3 features. After computing cosine similarity using

$$Sim(x,y) = \frac{x \cdot y}{||x||_2 \cdot ||y||_2}$$

For the bill subgraph, the role distribution is as shown in 5; for the campaign subgraph, it's shown in 6. The bill role distribution shows Legislator Node IDs 346 - Jerry Morgan of KS, 533 - Wicker Roger of MS, 1709 - Paul Simon of IL as top 3 most similar to the "average legislator"; while the campaign subgraph shows Node IDs 322 - Thomas "Tom" McClintock of CA, 1854 - Harold Washington of IL, 369 -Beto O'Rourke of TX as top 3 most similar to the "average legislator". None of them overlap. Clearly, the campaign subgraph's roles are not very meaningful as all nodes appear to have similar cosine similarities. We suspect that this is because we collated all the years together so that we lost data granularity in the process, and when one's been around for a while, he/she does the same thing for raising money - that is, he/she will take donations and build the money network.

Lastly, we wanted to understand how we can cut the graph efficiently, with the cut being a potential feature we can use later in link prediction in lieu of legislators' party allegiances. For the bill graph, we have the Positive Set of size and the Negative Set, $\overline{S}$ of size 933. For the campaign subgraph, we have the Positive Set, S of size 946 and the Negative Set, $\overline{S}$ of size 867, using the Clauset-Newman-Moore greedy modularity maximization provided by NetworkX. This closely resembles an even split of the aggregate two-party divide of the Congress.

Moreover, we recognize that so far, we are treating the 36 years' data as one aggregate graph - this probably aggravates the connectedness of the graph (as over time, one tends to collaborate with most people, and to get donations from all committees on the same of the aisle). Therefore, we also isolated one batch of the tripartite graph (defined as two years' campaign contribution data plus the two following years' bill co-authorship data). It turns out that even 2 years is enough time for the graph to become densely connected. We have plotted a few degree distribution plots for Campaign Year 1999-2000 with bill data from 2001 to 2002 (the 107th Congress) in Figures 7 and 8.

It holds true that the **Bill Co-authorship Graph resembles the academic collaboration graph with a power law pattern** (long tail) - the most frequent degrees are the smallest degrees, and it has a very high clustering coefficient. It's likely due to a few reasons: politics is a tiny field that everyone knows all the issues pretty well, and can have an opinion on almost anything and therefore removing the knowledge hurdle to co-sponsor bills; Congress is a two-party system, and legislators within one party tends to co-sponsor bills together along party lines; once more than two legislators sponsor a bill, they would create a triangle thus the high clustering coefficient; yet bills come in a dime a dozen, and it's unlikely that one legislator finds it necessary or efficient to co-sponsor every single bill he/she agrees with. Moreover, sponsoring bills together can signal an alliance, making legislators consider carefully before putting down their names.

For the **folded campaign contribution graph, it represents a typical small world pattern with high clustering coefficient**, when we look at legislators only, and a somewhat normal distribution - degrees are peaked between the smallest and the largest degrees. This makes sense as politics is a very small circle, and the national players and donors are relatively constant as they are mostly career politi-

Figure 7: Bill Subgraph Degree Distribution (log) for 2001-2002

Figure 8: Campaign Subgraph Degree Distribution (log) for 1999-2000

Figure 9: Selector Percentile for Feature Ranking

Figure 10: Grid Search for Optimal Tree Depth

cians.

4.2 Feature Selection

Using 100th and 101st Congress, we first observed that **features from contribution amounts** always rank dead last, while contributing to slow running time of our algorithms. Thus, we removed these features first (sum and absolute difference). We then tabulated results of running F-tests for the remaining features below:

Feature	F-Score
Clustering Coeff Difference	19265.8
Jaccard Index	14453.4
Degree Centrality Diff	10964.6
Shortest Distance	4615.9
Degree Difference	3438.7
Clustering Coeff Sum	1950.9
Common Neighbors	1391.9
Union of Neighbors	1215.6
Clustering Coeff Mean	841.1
If From Same Spectral Cluster	0.8

This indicates that knowing clusters generated from modularity maximization to mimic partyline is actually not helpful, which is a new learning to us. It also shows that the legislators' connectedness and the financial contribution communities they are in are important and indicative of their collaboration in office. The observation holds true when we re-ran the selection algorithm all the datasets available from the 98^{th} to the 112^{nd} Congress, with the same

Top 3 Features (see Apendix A).

In order to determine how many features we should be using in logistic regression, we used grid search to determine the optimal selector percentile as shown in Figure 9. This shows that we should be using all the features generated so far. For the decision tree, we have tuned the parameter for tree-depth, in order to avoid overfitting by running a grid search as shown in Figure 10. Therefore, we set the maximum tree depth to be 10.

4.3 Model Performance

We have run models in two ways: based on limited dataset (i.e. training on the 100^{th} Congress and test on the 101^{st} Congress, in what we call a one-term set), and based on richer datasets (i.e. training on the 98^{th} to 112^{nd} Congress combined graph, and test on the 113^{rd} to 114^{th} Congress combined graph). Below are our model performance in Tables 1 and 2, respectively . Candidate Only Predictor uses only the Affiliated Party and Home State information from candidates, gathered form the campaign contribution data.

To visualize models' accuracy in terms of true positives and true negatives, we have plotted confusion matrices for the limited dataset in Figures 11 and 12.

Model	Train Accuracy	Test Accuracy
Naive Baseline	0.695	0.695
Candidate Party/State	0.695	0.698
Logistic Regression	0.786	0.774
Decision Tree	0.854	0.794
Logistic Reg w/ Node2Vec	0.728	0.728

Table 1: Model Performance for Limited Dataset

Model	Train Accuracy	Test Accuracy
Naive Baseline	0.697	0.691
Candidate Party/State	0.695	0.698
Logistic Regression	0.748	0.714
Decision Tree	0.795	0.740

Table 2: Model Performance for All Datasets Combined

4.4 Discussion

It's interesting to note that knowing candidates' party and home state information does not lead to a better model when compared to the naive baseline, and using node2vec as we tried in this paper does not help either. Decision tree performed neck and shoulders above the rest, and a better performance in avoiding false positives and false negatives too when compared to logistic regression. This lends support to our hypothesis that money has a big influence in political collaboration: knowing the network structures of candidates' campaign donation graphs, we can reliably predict whom they wil collaborate with when elected, in the form of co-sponsoring bills.


Figure 11: Confusion Matrix for Decision Tree in Limited Dataset

Figure 12: Confusion Matrix for Logistic Regression in Limited Dataset

5 Conclusion

US Congressional Politics is indeed a small world: legislators are connected to other legislators via common donors and co-authorship on bills. We have identified the academic collaboration network-like pattern for bill co-authorship data, and a "small world" pattern among legislators, with consistently high clustering coefficients. Moreover, it does appear that "money moves politics": using features learned from campaign donation networks, we can confidently predict if two legislators will later collaborate on bills together - easily beating a naive baseline. In paricular, decision tree model performed very well to give us 82% accuracy.

This sheds new light on understanding politicians' behavior in Congress. Among many potential application, we now have a reliable way to predict if an elected candidate's campaign trail promises will likely carry through, by looking at whose money he/she has taken from, and with what other politicians does one share donors with.

6 Acknowledgement

References

[1] Mohammad Al Hasan, Vineet Chaoji, Saeed Salem, and Mohammed Zaki. Link predic-

tion using supervised learning. In *SDM06: workshop on link analysis, counter-terrorism and security*, 2006.

[2] Lars Backstrom and Jure Leskovec. Supervised random walks: predicting and recommending links in social networks. In *Proceedings of the fourth ACM international conference on Web search and data mining*, pages 635–644. ACM, 2011.

[3] Casey BK Dominguez. Groups and the party coalitions: A network analysis of overlapping donor lists. In *Paper delivered at the Annual Meeting of the American Political Science Association, Washington, DC*. Citeseer, 2005.

[4] James H Fowler. Connecting the congress: A study of cosponsorship networks. *Political Analysis*, 14(4):456–487, 2006.

[5] Aditya Grover and Jure Leskovec. node2vec: Scalable feature learning for networks. *Knowledge Discovery and Data Mining*, 2016.

[6] David Liben-Nowell and Jon Kleinberg. The link-prediction problem for social networks. *Journal of the American society for information science and technology*, 58(7):1019–1031, 2007.

[7] Rush Moody. Super-pac donor networks, Dec 2015.

[8] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

[9] Patricia Perozo, Gaspar Garcia, and Cheenar Banerjee. Co-authorship networks in congress, Dec 2016.

A   Feature Ranking with All Datasets

Feature	F–Score
Jaccard Index	180126.1
Degree Centrality Diff	99775.0
Clustering Coeff Difference	65108.8
Shortest Distance	49727.1
Common Neighbors	22229.7
Clustering Coeff Sum	19472.1
Union of Neighbors	9183.6
Degree Difference	8650.4
Clustering Coeff Mean	7026.5
Congress Term	3.6
If From Same Spectral Cluster	2.6